



SIEM01: Rebuilding the SIEM After a Hardware Failure

Purpose: Documents the recovery of the lab's log collection capability after the original Wazuh host failed, the hardware role swap that made the rebuild better than the original, and the Wazuh deployment on the new host. Written so a reviewer can see the reasoning behind the placement decision, not only the installation steps.

Host: SIEM01 , Dell OptiPlex 9020, Ubuntu 24.04.2 LTS, 192.168.20.2/24 on VLAN 20 (BlueTeam), switch Port 4.

Status: Wazuh 4.14.7 live since 2026-09-05. Three agents reporting as of 2026-09-06: DC01, ADM01 and PVE01.

1. What failed, and the decision that followed

The original SIEM was a physical machine running Rocky Linux 9.6 at 192.168.20.2 . It lost power on 2026-08-29 and did not return. That was the second hardware failure in this lab, after a server NIC, and it left the environment with no central log collection: pfSense continued logging locally, but nothing was being shipped or correlated.

The obvious response was to rebuild the same thing on the same kind of hardware. A better question was which machine *should* hold this role.

At that point the estate had two physical Linux hosts and a hypervisor with spare capacity. The monitoring host was a Dell OptiPlex 9020 with 8 cores and 16 GB of RAM, running Grafana and Prometheus, a workload that comfortably fits in about 2 GB. Meanwhile the service that needed rebuilding, Wazuh, ships an Indexer (an OpenSearch instance) whose practical memory floor is around 8 GB, and which accumulates data that must survive.

So the roles were the wrong way round, and the failure was the opportunity to correct it:

	Before	After
Dell OptiPlex 9020 (physical)	Grafana and Prometheus, VLAN 60	SIEM01 , Wazuh, VLAN 20
Proxmox guest	none	MON01 , Grafana and Prometheus, VLAN 60

The principle is worth stating plainly, because it generalises: **put the memory-hungry, data-bearing service on dedicated hardware, and the light, reproducible one on the hypervisor.** Grafana and Prometheus can be rebuilt from a configuration file in minutes. A SIEM's value is its accumulated history and the fact that it

keeps running while everything around it is being investigated.

```
nameservers:
  addresses: [8.8.8.8, 1.1.1.1, 8.8.4.4]
nantwi@nbl-core-ub01:~$ nproc && free -h && df -h /
8
              total        used        free      shared  buff/cache   available
Mem:           15Gi         856Mi         11Gi         1.4Mi         3.3Gi         14Gi
Swap:           4.0Gi           0B         4.0Gi
Filesystem                Size      Used Avail Use% Mounted on
/dev/mapper/ubuntu--vg-ubuntu--lv 455G      13G  424G   3% /
nantwi@nbl-core-ub01:~$
```

Figure 16.1: The candidate host: 8 cores, 15 GiB usable RAM, and 424 GB free of a 455 GB volume. The disk figure matters as much as the memory, because the Indexer is what eventually fills a SIEM's storage.

Full rationale and the naming decisions are in [docs/14](#) section 5.3.

2. Repurposing the host in place

The machine was **not** reinstalled. Ubuntu 24.04 carried over, which means Wazuh moved from Rocky Linux to Ubuntu. Three changes moved the host from one role to the other.

Renamed to match the convention ([docs/14](#) section 5.1):

```
sudo hostnamectl set-hostname SIEM01
sudo sed -i "s/nbl-core-ub01/SIEM01/g" /etc/hosts
```

The old name, `nbl-core-ub01`, encoded the operating system and nothing about the machine's purpose. That is precisely the failure mode the naming convention exists to prevent.

Readdressed from VLAN 60 to VLAN 20, and **recabled** from switch Port 8 to Port 4, the untagged VLAN 20 access port.

The cloud-init trap

The netplan file on an Ubuntu Server install is named `50-cloud-init.yaml`, and that name is a warning. It is generated by **cloud-init**, the program that configures a fresh machine on first boot. A file written by a program can be rewritten by that program, and a manual edit would be silently reverted at the next boot, leaving the host on its old address on a port that no longer carries that VLAN.

Disabling cloud-init's network management first is what makes the edit durable:

```
echo 'network: {config: disabled}' | sudo tee /etc/cloud/cloud.cfg.d/99-disable-network-config.cfg
```

The configuration then uses the current netplan syntax. `gateway4:` is deprecated in 24.04 and replaced by a `routes:` entry, and DNS points at pfSense rather than a public resolver, matching DC01 and KALI01:

```
network:
  version: 2
```

```
ethernets:
  eno1:
    dhcp4: no
    addresses:
      - 192.168.20.2/24
    routes:
      - to: default
        via: 192.168.20.1
    nameservers:
      addresses: [192.168.20.1]
```

```
nantwi@SIEM01:~: command not found
nantwi@SIEM01:~$ sudo cat /etc/netplan/50-cloud-init.yaml
network:
  version: 2
  ethernets:
    eno1:
      dhcp4: no
      addresses:
        - 192.168.20.2/24
      routes:
        - to: default
          via: 192.168.20.1
      nameservers:
        addresses: [192.168.20.1]
nantwi@SIEM01:~$ sudo netplan generate && echo "config OK"
config OK
```

Figure 16.2: The VLAN 20 configuration after a reboot, with `netplan generate` confirming it parses. That it survived the reboot is the evidence that cloud-init is genuinely no longer managing the interface.

The general lesson: on Linux, a configuration file named after a program is usually owned by that program. Editing it works until the program runs again. Either turn the program off, as here, or place your settings in a higher-numbered file that wins.

3. Two failures worth recording

Applying the network change before moving the cable. `netplan apply` was run while the machine was still plugged into the VLAN 60 port. The host immediately held an address that did not exist on the network it was attached to: unreachable, with no route to its gateway. Nothing was damaged, and moving the cable resolved it, but the correct order is to write the configuration, shut down, move the cable, then power on. A machine should never be left holding an address its physical connection cannot serve.

SSH refused the connection afterwards. The address `192.168.20.2` had belonged to the failed Rocky Linux host, whose SSH host key was still recorded in the administrator's `known_hosts`. A different physical machine now answered at that address with a different key, so SSH refused:

```
WARNING: REMOTE HOST IDENTIFICATION HAS CHANGED!
```


This is the control working exactly as designed. SSH cannot distinguish "the hardware was replaced" from "something is intercepting this connection", so it stops and requires a human decision. The response was to verify rather than dismiss: the fingerprint was read directly from the host itself,

```
sudo ssh-keygen -lf /etc/ssh/ssh_host_ed25519_key.pub
```

confirmed identical to the one SSH had presented, and only then was the stale record removed with `ssh-keygen -R 192.168.20.2`.

SSH also volunteered corroboration, noting the same key was already known under `192.168.60.2`, the machine's previous address. That is independent confirmation that this was a host that moved, not a stranger.

The distinction that matters: deleting the warning and reconnecting is what most people do, and it works. Verifying the fingerprint out of band is what turns a dismissed alert into a cleared one. If that warning ever appears when nothing has changed, it is the one time it is telling you something real.

A terminal window titled 'Windows PowerShell' with a tab for 'nantwi@SIEM01: ~'. The user enters a complex command: `ip -br a && ip r && ping -c2 192.168.20.1 && nslookup github.com`. The output shows network details for interfaces `lo`, `eno1`, and `docker0`, including IP addresses and link states. It then shows the results of a ping to `192.168.20.1` and an `nslookup` for `github.com`.


```
nantwi@SIEM01:~$ ip -br a && ip r && ping -c2 192.168.20.1 && nslookup github.com
lo                UNKNOWN      127.0.0.1/8  ::1/128
eno1              UP          192.168.20.2/24 fe80::6600:6aff:fe5b:bd63/64
docker0          DOWN        172.17.0.1/16
default via 192.168.20.1 dev eno1 proto static
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1 linkdown
192.168.20.0/24 dev eno1 proto kernel scope link src 192.168.20.2
PING 192.168.20.1 (192.168.20.1) 56(84) bytes of data:
64 bytes from 192.168.20.1: icmp_seq=1 ttl=64 time=1.44 ms
64 bytes from 192.168.20.1: icmp_seq=2 ttl=64 time=0.432 ms

--- 192.168.20.1 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1002ms
rtt min/avg/max/mdev = 0.432/0.936/1.441/0.504 ms
Server:      127.0.0.53
Address:     127.0.0.53#53

Non-authoritative answer:
Name:   github.com
Address: 140.82.112.3

nantwi@SIEM01:~$
```

Figure 16.3: Post-reboot verification: the static address held, the default route is marked `proto static`, the gateway answers, and name resolution works. Four separate things confirmed in one command.

A terminal window showing the user running `192.168.20.1 ping statistic` and `resolvectl status eno1 | grep -A3 "DNS Servers"`. The output shows the upstream resolver is `192.168.20.1`.

```
nantwi@SIEM01:~$ 192.168.20.1 ping statistic
192.168.20.1: command not found
nantwi@SIEM01:~$ resolvectl status eno1 | grep -A3 "DNS Servers"
    DNS Servers: 192.168.20.1
nantwi@SIEM01:~$
```

Figure 16.4: `resolvectl` confirming the upstream resolver is `192.168.20.1`, pfSense, rather than a leftover public server. `127.0.0.53` in an `nslookup` result is systemd-resolved, a local stub that forwards on; it is worth checking what sits behind it.

4. Firewall consequences

Two rules were required before the host could be administered, and both were built before the move rather than after.

`MGMT-11` permits the management plane to administer the SIEM, using aliases so the rule reads by intent rather than by number:

Alias	Type	Value
<code>SIEM01_HOST</code>	Host	<code>192.168.20.2</code>
<code>SIEM_ADMIN</code>	Ports	22 (SSH), 443 (Wazuh dashboard)

Port 443 was included before Wazuh existed, on the expectation that the dashboard would use it. The installer confirmed this during the run (`wazuh web interface port will be 443`), so no rule change was needed afterwards.

The rule is **logged**, on the same reasoning as `MGMT-08` on the domain controller: administrative access to a machine holding security evidence is privileged activity and has to be attributable (NIST AC-17, AU-2). Registered in `docs/13` as `MGMT-11`.

The screenshot shows the 'Firewall / Aliases / Edit' interface. The 'Properties' section includes: Name (SIEM01_HOST), Description (SIEM01 Wazuh server, VLAN 20 BlueTeam), and Type (Host(s)). The 'Host(s)' section includes a Hint and two input fields: 'IP or FQDN' (192.168.20.2) and 'SIEM01'. At the bottom are buttons for 'Save', 'Export to file', and 'Add Host'.

Figure 16.5: The `SIEM01_HOST` alias. Defining the host once means the rule references intent rather than an address, and re-addressing the machine later becomes a single edit instead of a hunt through the ruleset.

Firewall / Aliases / Edit

Properties

Name
The name of the alias may only consist of the characters "a-z, A-Z, 0-9 and _".

Description
A description may be entered here for administrative reference (not parsed).

Type

Port(s)

Hint Enter ports as desired, with a single port or port range per entry. Port ranges can be expressed by separating with a colon.

Port	Description	Action
22	SSH	Delete
443	Wazuh dashboard	Delete

Figure 16.6: The **SIEM_ADMIN** ports alias, holding SSH and the dashboard with each port individually described. Port 443 was added before Wazuh was installed, on the expectation that the dashboard would use it.

Rule	Status	Name	Protocol	Action	Description	Logging
MGMT-02	Enabled	DNS resolution via pfSense resolver. NIST SC-7	53 (DNS)	none	MGMT-02 - DNS resolution via pfSense resolver. NIST SC-7	☒
MGMT-03	Enabled	NTP time sync from pfSense; correct time underpins logging and Kerberos. NIST AU-8	123 (NTP)	none	MGMT-03 - NTP time sync from pfSense; correct time underpins logging and Kerberos. NIST AU-8	☒
MGMT-04	Enabled	ICMP echo for troubleshooting reachability. NIST CA-7	*	none	MGMT-04 - ICMP echo for troubleshooting reachability. NIST CA-7	☒
MGMT-05	Enabled	AD domain membership (TCP) to SRV1 DC. NIST AC-4	AD_TCP	none	MGMT-05 - AD domain membership (TCP) to SRV1 DC. NIST AC-4	☒
MGMT-06	Enabled	AD domain membership (UDP) to SRV1 DC. NIST AC-4	AD_UDP	none	MGMT-06 - AD domain membership (UDP) to SRV1 DC. NIST AC-4	☒
MGMT-07	Enabled	RPC dynamic ports (domain join, GPO, replication) to SRV1 DC. NIST AC-4	AD_RPC_DYNAMIC	none	MGMT-07 - RPC dynamic ports (domain join, GPO, replication) to SRV1 DC. NIST AC-4	☒
MGMT-08	Enabled	Admin access (RDP + WinRM/Ansible) to SRV1 DC. NIST AC-6/AC-17/AU-2	MGMT_TCP	none	MGMT-08 - Admin access (RDP + WinRM/Ansible) to SRV1 DC. NIST AC-6/AC-17/AU-2	☒
MGMT-09	Enabled	Admin reach into all lab VLANs. NIST AC-6	LAB_NETS	none	MGMT-09 - Admin reach into all lab VLANs. NIST AC-6	☒
MGMT-10	Enabled	Internet access (anywhere except the lab). NIST SC-7	! LAB_NETS	none	MGMT-10 - Internet access (anywhere except the lab). NIST SC-7	☒
MGMT-11	Enabled	SSH and dashboard access to SIEM01 for administration. NIST AC-17, AU-2	SIEM01_HOST	SIEM_ADMIN	MGMT-11 - SSH and dashboard access to SIEM01 for administration. NIST AC-17, AU-2	☒
All	Enabled	All	*	none	All	☒

Figure 16.7: **MGMT-11** at the foot of the **MANAGEMENT** tab, both aliases resolving as links and the logging icon set. The rules above it show the rest of the management plane's ruleset carrying live traffic.

The DNS Resolver check. Before moving the host, the pfSense DNS Resolver was confirmed to be listening on the BLUETEAM interface. This is the third time in this build that a service bound to the wrong interfaces would have produced a correct-looking configuration and a silent failure, after the same issue on VLAN 30 (docs/15 section 6). It is now a standing pre-flight check when a host moves to a new segment.

5. Installing Wazuh

Wazuh is three programs, and knowing which is which is the difference between diagnosing a problem and guessing at one:

Component	Job	Listens on
Wazuh Indexer	Stores and searches events. OpenSearch underneath. The memory-hungry component.	9200 (localhost)
Wazuh Server	The manager. Agents connect to it; it applies detection rules and forwards results.	1514, 1515, 55000
Wazuh Dashboard	The web interface.	443
(Filebeat)	Ships the manager's own alerts into the Indexer.	n/a

"Wazuh is down" is never a useful diagnosis. It is one of these.

The all-in-one installer deploys all three on a single host, which suits 16 GB and 8 cores. A distributed deployment only earns its complexity when the Indexer outgrows one machine.

The agent and manager conflict

The first install attempt was refused:

```
ERROR: Wazuh manager already installed.
```

Investigation showed no manager, only a `wazuh-agent` left from when this machine was the monitoring host reporting to the old SIEM.

A manager and an agent cannot coexist on one host, because both packages own `/var/ossec`. It is not a limitation to work around: the manager monitors its own host natively, so an agent on the manager would be redundant as well as conflicting.

```
nantwi@SIEM01:~$ 192.168.20.1 ping statistic
192.168.20.1: command not found
nantwi@SIEM01:~$ resolvectl status eno1 | grep -A3 "DNS Servers"
DNS Servers: 192.168.20.1
nantwi@SIEM01:~$ curl -sO https://packages.wazuh.com/4.14/wazuh-install.sh && sudo bash ./wazuh-install.sh -a
[sudo] password for nantwi:
05/09/2026 19:51:26 INFO: Starting Wazuh installation assistant. Wazuh version: 4.14.7
05/09/2026 19:51:26 INFO: Verbose logging redirected to /var/log/wazuh-install.log
05/09/2026 19:51:27 ERROR: Wazuh manager already installed.
05/09/2026 19:51:27 INFO: If you want to overwrite the current installation, run this script adding the option -o/--overwrite
. This will erase all the existing configuration and data.
nantwi@SIEM01:~$ ls
bluelab-ansible Docker-Lessons test_winrm.py 'udo systemctl restart grafana-server' wazuh-install.sh
nantwi@SIEM01:~$ sudo bash ./wazuh-install.sh -a
05/09/2026 19:52:02 INFO: Starting Wazuh installation assistant. Wazuh version: 4.14.7
05/09/2026 19:52:02 INFO: Verbose logging redirected to /var/log/wazuh-install.log
05/09/2026 19:52:03 ERROR: Wazuh manager already installed.
05/09/2026 19:52:03 INFO: If you want to overwrite the current installation, run this script adding the option -o/--overwrite
. This will erase all the existing configuration and data.
nantwi@SIEM01:~$
```

Figure 16.8: The installer refusing, and the investigation that followed. Only `wazuh-agent 4.11.2` was present, a leftover from the machine's previous role.

The installer offers `-o/--overwrite`, which erases all existing Wazuh configuration and data. Removing the single package explicitly was preferred, because the overwrite flag is a blunt instrument and on a machine whose job is to be trustworthy it is worth knowing exactly what changed:

```
sudo systemctl disable --now wazuh-agent
sudo apt purge -y wazuh-agent
sudo rm -rf /var/ossec

nantiw@SIEM01:~$ sudo systemctl disable --now wazuh-agent && sudo apt purge -y wazuh-agent && sudo rm -rf /var/ossec
Removed "/etc/systemd/system/multi-user.target.wants/wazuh-agent.service".
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages will be REMOVED:
  wazuh-agent*
0 upgraded, 0 newly installed, 1 to remove and 0 not upgraded.
After this operation, 39.7 MB disk space will be freed.
(Reading database ... 133547 files and directories currently installed.)
Removing wazuh-agent (4.11.2-1) ...
(Reading database ... 133141 files and directories currently installed.)
Purging configuration files for wazuh-agent (4.11.2-1) ...
nantiw@SIEM01:~$ dpkg -l | grep -i wazuh; ls /var/ossec 2>&1; echo "--- clean if both empty ---"
ls: cannot access '/var/ossec': No such file or directory
--- clean if both empty ---
nantiw@SIEM01:~$ free -h && sudo ss -tlnp | grep -E ':(443|1514|1515|9200|55000)\s'; echo "--- ports clear if nothing above ---"
              total        used         free       shared    buff/cache   available
Mem:           15Gi        624Mi        14Gi         1.4Mi         930Mi        14Gi
Swap:          4.0Gi           0B         4.0Gi
--- ports clear if nothing above ---
nantiw@SIEM01:~$ sudo bash ./wazuh-install.sh -a
05/09/2026 19:58:03 INFO: Starting Wazuh installation assistant. Wazuh version: 4.14.7
05/09/2026 19:58:03 INFO: Verbose logging redirected to /var/log/wazuh-install.log
05/09/2026 19:58:12 INFO: Verifying that your system meets the recommended minimum hardware requirements.
05/09/2026 19:58:12 INFO: Wazuh web interface port will be 443.
05/09/2026 19:58:16 INFO: --- Dependencies ---
05/09/2026 19:58:16 INFO: Installing debhelper.
```

Figure 16.9: The agent purged, `/var/ossec` gone, Wazuh's ports confirmed free, and the installer proceeding past the check that had blocked it.

Running the install

The installer was downloaded to a file rather than piped into a shell:

```
curl -sO https://packages.wazuh.com/4.14/wazuh-install.sh
sudo bash ./wazuh-install.sh -a
```

The pattern `curl ... | sudo bash` is common and is worth avoiding on principle. It executes code as root that was never seen, from a server that could have been compromised between publication and retrieval. Downloading first costs nothing and makes inspection possible.


```

05/09/2026 19:59:46 INFO: wazuh-indexer service started.
05/09/2026 19:59:46 INFO: Initializing Wazuh indexer cluster security settings.
05/09/2026 19:59:52 INFO: Wazuh indexer cluster security configuration initialized.
05/09/2026 19:59:52 INFO: Wazuh indexer cluster initialized.
05/09/2026 19:59:52 INFO: --- Wazuh server ---
05/09/2026 19:59:52 INFO: Starting the Wazuh manager installation.
05/09/2026 20:01:08 INFO: Wazuh manager installation finished.
05/09/2026 20:01:09 INFO: Wazuh manager vulnerability detection configuration finished.
05/09/2026 20:01:09 INFO: Starting service wazuh-manager.
05/09/2026 20:01:25 INFO: wazuh-manager service started.
05/09/2026 20:01:25 INFO: Starting Filebeat installation.
05/09/2026 20:01:48 INFO: Filebeat installation finished.
05/09/2026 20:01:52 INFO: Filebeat post-install configuration finished.
05/09/2026 20:01:52 INFO: Starting service filebeat.
05/09/2026 20:01:53 INFO: filebeat service started.
05/09/2026 20:01:53 INFO: --- Wazuh dashboard ---
05/09/2026 20:01:53 INFO: Starting Wazuh dashboard installation.
05/09/2026 20:10:38 INFO: Wazuh dashboard installation finished.
05/09/2026 20:10:38 INFO: Wazuh dashboard post-install configuration finished.
05/09/2026 20:10:38 INFO: Starting service wazuh-dashboard.
05/09/2026 20:10:38 INFO: wazuh-dashboard service started.
05/09/2026 20:10:40 INFO: Updating the internal users.
05/09/2026 20:10:45 INFO: A backup of the internal users has been saved in the /etc/wazuh-indexer/internalusers-backup folder.
05/09/2026 20:10:54 INFO: The filebeat.yml file has been updated to use the Filebeat Keystore username and password.
05/09/2026 20:11:25 INFO: Initializing Wazuh dashboard web application.
05/09/2026 20:11:27 INFO: Wazuh dashboard web application initialized.
05/09/2026 20:11:27 INFO: --- Summary ---
05/09/2026 20:11:27 INFO: You can access the web interface https://<wazuh-dashboard-ip>:443

```

Figure 16.10: The Indexer, manager, Filebeat and Dashboard each installing and starting in sequence. This screenshot is deliberately cropped: the installer prints the generated admin credentials immediately below, and they must not reach a repository.




Figure 16.11: The dashboard on first login. "No agents registered" alongside 336 alerts is not a contradiction: the manager monitors its own host as agent `000`.

6. Credentials, and why the dashboard could not change them

The installer prints a generated `admin` password once. Attempting to change it through the dashboard fails:

```
Failed to reset password. {"status":"FORBIDDEN","message":"Resource 'admin' is reserved."}
```




Figure 16.12: The dashboard refusing to change its own administrative password. The account is a reserved resource in the Indexer's security configuration and cannot be modified through the API the interface uses.

This is a deliberate boundary rather than a defect. The `admin` account is defined in the security plugin's `internal_users.yml` and marked **reserved**, which makes it immutable from the application layer. A compromised dashboard session therefore cannot lock the owner out of their own cluster. Changing it requires filesystem access to the host, which is a meaningfully higher bar:

```
curl -sO https://packages.wazuh.com/4.14/wazuh-passwords-tool.sh
sudo bash wazuh-passwords-tool.sh -u admin -p '<new password>'
```

The tool rewrites the hash and reruns `securityadmin` to push the change into the cluster.

An evidence-handling rule came out of this. The moment an installer prints a generated credential is exactly the moment people screenshot it. Figure 16.10 is cropped for that reason. Every screenshot is now checked for passwords, keys and tokens before it enters `images/`, and a credential that has appeared in a screenshot is rotated rather than trusted.

7. Post-install configuration

Removing the previous role's software

Grafana and Prometheus were left running through the Wazuh install so nothing was disturbed mid-deployment, then removed. A SIEM holds the evidence for everything else in the environment, which makes unrelated software on it attack surface with no purpose (NIST CM-7).

Grafana came from an apt repository and removed cleanly. Prometheus did not, and the reason is worth knowing:

```
sudo apt purge -y grafana prometheus prometheus-node-exporter
```

```
Package 'prometheus' is not installed, so not removed
```

while `systemctl is-active prometheus` still reported **active**. Prometheus is commonly installed from the project's own tarball rather than a distribution package, so apt has no record of it even though it is plainly running:

```
which prometheus promtool
ls -d /etc/prometheus /var/lib/prometheus /etc/systemd/system/prometheus.service
```

Binaries in `/usr/local/bin`, a hand-written unit file, and data directories. Removing it means removing those directly:

```
sudo systemctl disable --now prometheus
sudo rm -f /etc/systemd/system/prometheus.service
sudo systemctl daemon-reload
sudo rm -rf /usr/local/bin/prometheus /usr/local/bin/promtool /etc/prometheus /var/lib/prometheus
sudo userdel prometheus
```

The general point: "apt says it is not installed" is not the same as "it is not on the machine". `/usr/local` exists precisely for software the package manager does not own, and asking systemd and the filesystem is how you find it. On a host being repurposed, check both.

Dashboard session lifetime, and an outage caused by fixing it

The session expired quickly enough to interrupt work. Two settings in `/etc/wazuh-dashboard/opensearch_dashboards.yml` govern it, and the distinction between them is the useful part:

Setting	What it controls
<code>opensearch_security.session.ttl</code>	How long a session may live, in milliseconds
<code>opensearch_security.session.keepalive</code>	Whether that clock restarts on each interaction

Without `keepalive` a session ends a fixed time after login regardless of activity, so it can expire mid-investigation. With it, the timeout becomes an **idle** timeout.

The first attempt took the dashboard down. The change was appended to the end of the file. The service then failed to start, and `systemctl status` reported repeated restart attempts until systemd gave up:

```
wazuh-dashboard.service: Start request repeated too quickly.  
wazuh-dashboard.service: Failed with result 'exit-code'.
```

Recovery was immediate because a backup had been taken in the same breath as the edit:

```
sudo cp /etc/wazuh-dashboard/opensearch_dashboards.yml.bak /etc/wazuh-dashboard/opensearch_dashboa  
sudo systemctl reset-failed wazuh-dashboard  
sudo systemctl start wazuh-dashboard
```

`reset-failed` is required: systemd rate-limits a service that keeps failing, and until that counter is cleared `start` silently does nothing.

The cause was duplicate keys. One `grep` found it:

```
17:opensearch_security.cookie.ttl: 900000  
18:opensearch_security.session.ttl: 900000  
19:opensearch_security.session.keepalive: true
```

The settings were already present. Appending them defined each key twice, and YAML rejects a document that does that, so the configuration would not parse. The defaults also explained the original complaint exactly: `900000` ms is **15 minutes**, and `keepalive` was already enabled, which is why the session survived activity and expired during idle time.

The correct change edits the existing lines rather than adding new ones:

```
sudo sed -i 's/^opensearch_security.cookie.ttl: .*/opensearch_security.cookie.ttl: 3600000/; s/^op  
sudo systemctl restart wazuh-dashboard
```

Three lessons, in order of usefulness. Copy a configuration file before editing it, since the backup turned a failed service into a thirty-second recovery. Check whether a setting already exists before adding it, because appending to a configuration file is not the same as configuring it. And editing a file changes what a service *will* load, not what it is currently running, so the restart and the verification are part of the change rather than optional extras.

One hour of inactivity was chosen deliberately rather than for convenience. A long-lived authenticated session on the machine holding all the security evidence is itself a risk, since an unattended workstation stays logged in. A regulated environment would more likely keep the shipped fifteen minutes, backed by single

sign-on and MFA, which is the direction the companion IAM lab is building toward. One hour is the lab compromise and is recorded here so the decision is visible rather than assumed (NIST AC-11, Session Lock; AC-12, Session Termination).

A dormant privileged account

Authenticating a systemd action produced a polkit prompt offering two identities:

1. Noble Antwi (nantwi)
2. Ansible Automation Service Account (ansible)

The second is a local account created when this machine was the monitoring host managed by the Ansible controller. `groups ansible` returns `ansible sudo`, so it holds full administrative rights, and `ANS01` no longer exists, so nothing uses it.

A dormant account with sudo on the host that stores the security evidence is exactly the finding an access review exists to catch. It was locked rather than deleted, because the controller is being rebuilt and will need an account:

```
sudo passwd -l ansible
```

`passwd -l` disables password authentication but does **not** prevent key-based SSH, so this is a partial control. It is recorded as `H-03` in the hardening register in `docs/13` for review when `ANS01` returns, rather than left to drift (NIST AC-2, Account Management; CA-5).

8. What it monitors today

With no agents deployed, the dashboard already shows several hundred alerts. They all carry agent ID `000`, which is always the manager itself.

Most are **Security Configuration Assessment** findings: Wazuh ships CIS Benchmark policies and evaluates the host against them at startup. These are not detections of an attack. They are an audit of the machine's own hardening, and they are useful for two reasons.

The first is direct: a SIEM holds evidence of everything else in the environment, which makes it a high-value target, and the tool has just produced a prioritised list of its own weaknesses.

The second is conceptual. SCA is continuous configuration monitoring, which is the same control family being satisfied by hand for pfSense in `docs/13`: NIST CM-2 (Baseline Configuration) and CM-6 (Configuration Settings). The firewall register is a manual implementation of what SCA does automatically for hosts.

The baseline, recorded before any hardening

Filtering the Configuration Assessment events by result gives the starting position against the **CIS Ubuntu Linux 24.04 LTS Benchmark v1.0.0**:

Result	Checks
Passed	147
Failed	127
Total evaluated	274
Pass rate	53.6%




Figure 16.13: 127 failing checks at first assessment. The visible failures are representative of a stock installation: `auditd` not installed, AIDE absent, no remote log host configured, audit tool ownership unset.




Figure 16.14: 147 passing checks. Most are account and permission hygiene that Ubuntu gets right by default, such as shadowed passwords and no duplicate UIDs.

Roughly half a benchmark failing is normal for a default installation, and that is the point: hardening is work that has to be done, not a property an operating system arrives with.

This baseline cannot be recreated later, which is why it was captured before anything was changed. It is the "before" half of a before-and-after, in the same shape as the management-plane isolation test in [docs/13](#) section 3.1, where a control was measured, changed, and measured again.

Planned as an exercise: harden a set of the failing checks, starting with the audit subsystem, and show the pass rate improving against these numbers.

A note on the Inventory tab. It is empty, and that is expected rather than a fault. Configuration Assessment's Inventory view is scoped to an agent selected from the agents list, and the manager is agent `000`, which does not appear there. The Events tab reads the alerts index directly, which is why the manager's own results are visible at all. Once agents enrol, their assessments populate the Inventory view normally.

9. The first agent: DC01

Deploying the first agent is where the interface-direction reasoning stops being theory.

Agents dial out to the manager. The manager listens. So the rule that lets DC01 report does not belong on the BlueTeam tab; it belongs on **ENTERPRISELAN**, the interface DC01's traffic arrives on. Registered as `ENT-05` in [docs/13](#).

Installing it

The deployment command comes from the dashboard (**Agents** → **Deploy new agent**), which fills in the manager address and matches the agent version to the manager. Run in PowerShell as Administrator on DC01:

```
Invoke-WebRequest -Uri https://packages.wazuh.com/4.x/windows/wazuh-agent-4.14.7-1.msi -OutFile $env:TEMP\wazuh-agent.msi /q WAZUH_MANAGER='192.168.20.2' WAZUH_REGISTRATION_SERVER=192.168.20.2 NET START WazuhSvc
```

Verification from the agent side, which is the diagnostic worth running first when an agent does not appear:

```
Get-Service WazuhSvc
Test-NetConnection 192.168.20.2 -Port 1514
```

It did not connect, and why

The agent installed and the service started, but nothing appeared in the dashboard. Two causes, in sequence.

The first was the rule itself, described in [docs/13](#) section 2.2: `ENT-05` had been copied from `MGMT-11` and carried the wrong source and service, so it matched nothing.

The second, after the rule was corrected, was that **the filter had not been reloaded**. The rule was right, saved, and applied, and the agent still could not connect until **Status** → **Filter Reload**.

That is the third time a stale filter has produced a different-looking failure in this lab. It has now been promoted to principle 9 of the rulebase in [docs/13](#): a saved rule is not an enforced rule.

The general diagnostic. An agent that cannot reach its manager is indistinguishable, from the dashboard, from an agent that is broken. Both are simply absent. `Test-NetConnection <manager> -Port 1514` from the endpoint separates the two in seconds: if the port test fails, the problem is the network path and nothing on the endpoint needs touching.

The result




Figure 16.15: Agent 001, DC01, active, Windows Server 2025, reporting from 192.168.50.2. System inventory, MITRE ATT&CK tactic counts, vulnerability detection and a CIS benchmark run all populate from a single agent install.




Figure 16.16: One agent active, none disconnected, and the alert counts on the overview now include a real endpoint rather than the manager alone.

Recording: Deploying a Wazuh Agent to Windows Server 2025, and debugging the pfSense rule that blocked it (5:34). The full deployment including both failures and the diagnosis, rather than an edited clean run. The raw file is not committed; see `images/README.md` for the recording policy.

The second agent, and what it revealed about the network

The administrative workstation was enrolled next as `ADM01`, using the same procedure. One rule covers it and every future management-plane agent:

```
MGMT-12 : source MANAGEMENT subnets , destination SIEM01_HOST , ports WAZUH_AGENT .
```

The interesting part is what the SIEM recorded. The agent registered with the IP address `100.118.195.0`, which is a Tailscale address, not the laptop's VLAN 10 address of `192.168.10.3`.

The traffic had not crossed the firewall at all. pfSense acts as a Tailscale subnet router and advertises the lab networks into the tailnet, so the laptop holds a route saying `192.168.20.0/24` is reachable over the Tailscale interface. Windows selects the most specific matching route, and a `/24` beats the default route through `192.168.10.1`. The overlay therefore won, even with the laptop sitting on VLAN 10 a few metres from the firewall.

`MGMT-12` was correct, applied, and had passed no traffic. The SIEM's own asset inventory is what exposed it.

Forcing the LAN path is a matter of beating specificity with more specificity, since a host route beats a `/24`:

```
route -p add 192.168.20.2 mask 255.255.255.255 192.168.10.1 metric 1
Restart-Service WazuhSvc
```

After reconnection the agent re-registered as `192.168.10.3`, which is the evidence that `MGMT-12` is genuinely carrying traffic.

The route degrades sensibly. Away from home, `192.168.10.1` is unreachable, the route leaves the active table, and the agent falls back to Tailscale. On the LAN it uses the LAN; roaming, it uses the overlay.




Figure 16.17: DC01 and ADM01 both active. ADM01's registered address is `192.168.10.3` rather than its earlier tailnet address, which is what confirms the management-plane rule is being exercised.

Two things worth carrying forward. An agent's registered address is whichever source address the manager observed, so it reports the path taken rather than the machine's primary identity. And a firewall rule that has never passed traffic proves nothing, however correct it looks: this is the same lesson as the stale filter, arriving by a different route.

The third agent: the hypervisor, and three platform surprises

`PVE01`, the Proxmox host, was enrolled last and is arguably the most important of the three. It runs KALI01 today and will run MON01 and VAULT01. Compromise it and every guest goes with it, which makes it the highest-value host in the lab after the domain controller. Nothing had been watching it.

It needed no new firewall rule. `MGMT-12` was written against the management subnet rather than a single address, so the hypervisor at `192.168.10.6` was already covered.

The route was checked before installing anything, which is the habit the ADM01 episode produced:

```
ip route get 192.168.20.2
```

```
192.168.20.2 via 192.168.10.1 dev vmbr0.10 src 192.168.10.6
```

Through pfSense, out of the VLAN 10 interface, from the host's real address. No Tailscale client on this machine, so no overlay route could win. This agent's traffic genuinely crosses the firewall, which makes it the honest validation of `MGMT-12`.

The install then failed three times, each for a different reason, and each worth recording because they are the same class of problem: **a vendor's generated command assumes a platform you may not be running.**

sudo: command not found. The dashboard generates `sudo dpkg -i ...`. Proxmox is a minimal Debian where you operate as root, and `sudo` is not installed. The `wget` half of the chained command succeeded, `sudo` failed, and `dpkg` never ran, so the service did not exist. The download looking successful made it appear the install had worked.

wazuh-agent depends on lsb-release; however: Package lsb-release is not installed. `dpkg -i` installs exactly the file it is given and does not resolve dependencies; that is `apt`'s job. Most Debian systems carry `lsb-release`; a minimal Proxmox install does not. The package unpacked but was left unconfigured, so the service still failed to start.

Invalid server address found: 'MANAGER_IP'. This is the instructive one. After installing the dependency and re-running, the agent started and immediately exited:

```
wazuh-agentd: ERROR: (4112): Invalid server address found: 'MANAGER_IP'
wazuh-agentd: ERROR: (1215): No client configured. Exiting.
```

`MANAGER_IP` is the literal placeholder in Wazuh's shipped configuration template. The `WAZUH_MANAGER` and `WAZUH_AGENT_NAME` environment variables are read by the package's configure script **on first install only**. Because the earlier attempt had already unpacked the package, the re-run was treated as a reinstall, the substitution was skipped, and the placeholder survived. The variables were on the command line and had no effect.

Repeating the install cannot fix this. A purge can:

```
systemctl stop wazuh-agent 2>/dev/null; apt-get purge -y wazuh-agent && rm -rf /var/ossec  
WAZUH_MANAGER='192.168.20.2' WAZUH_AGENT_GROUP='default' WAZUH_AGENT_NAME='PVE01' dpkg -i ./wazuh-
```

And the configuration is verified **before** starting the service, rather than after it fails:

```
grep -A3 "<server>" /var/ossec/etc/ossec.conf
```

```
<server>  
  <address>192.168.20.2</address>  
  <port>1514</port>  
  <protocol>tcp</protocol>
```

The general lesson. A failed installation is not a neutral state. It leaves artifacts that change how the next attempt behaves, and the second attempt can fail in a way that has nothing to do with the original cause. When an install fails partway, purge before retrying rather than running the same command again.




Figure 16.18: Three agents active across two operating system families: two Windows and one Debian.




Figure 16.19: The hypervisor as the SIEM now sees it. Hardware inventory, a CIS Debian 13 benchmark result, and 313 detected vulnerabilities including 22 rated critical, on a machine that had no monitoring at all an hour earlier.

10. Four baselines, and what the comparison actually shows

Every agent runs its platform's CIS benchmark on enrolment, and the manager assesses itself. Four hosts, four operating systems, captured before any hardening:




Figure 16.20: The four baselines. Bar length is proportional to the number of checks in each benchmark rather than normalised, because the size difference is most of the explanation for the score difference.

Host	Role	Benchmark	Passed	Failed	N/A	Checks	Score
DC01	Domain controller	CIS Microsoft Windows Server 2025	105	293	0	398	26%
ADM01	Administrative workstation	CIS Microsoft Windows 11 Enterprise v3.0.0	124	348	10	482	26%
PVE01	Hypervisor	CIS Debian Linux 13	80	105	22	207	43%
SIEM01	SIEM	CIS Ubuntu Linux 24.04 LTS	147	127	n/a	274	53.6%

SIEM01's figures come from assessment events rather than a policy summary, because the Wazuh manager is agent `000` and does not appear in the agent list, so the Configuration Assessment view cannot be opened for it. The arithmetic is identical: passed divided by passed plus failed.

The obvious reading is the wrong one

This is **not** a ranking of how secure the four operating systems are, and presenting it as one would be a mistake.

The **number of checks** is the missing variable. The Windows benchmarks contain roughly twice as many checks as the Linux ones: 482 and 398 against 274 and 207. They cover hundreds of Group Policy settings, audit subcategories, user rights assignments and security options that have no equivalent concept on a Linux host. A larger and more prescriptive benchmark produces more failures against a default installation, so the score falls.

That is why the chart above keeps the bars proportional to benchmark size instead of normalising each to 100%. Normalising would make the four look directly comparable, which is precisely the misreading to avoid.

What the numbers are actually for

Measuring a host against its own future. DC01 at 26% today, against DC01 in a month, is a meaningful statement about work performed. DC01 at 26% against SIEM01 at 53.6% is not a statement about anything.

Two conclusions do carry across all four, and they are the useful ones:

Every default installation fails most of its benchmark. Roughly half on Linux, three quarters on Windows. Hardening is deliberate work that has to be scheduled and done; it is not a property an operating system arrives with. Four independent data points make that hard to argue with.

The most important host scores worst. DC01 is the machine every other machine trusts: compromise it and every authentication decision in the environment is compromised with it. It has one of the largest benchmarks and the lowest score. ADM01 sits alongside it at 26%, and that pairing is uncomfortable for a reason, since the administrative workstation is where domain credentials are typed. Those two are the obvious place to start, and they are the two hosts already carrying `H-01` in the hardening register.




Figure 16.21: DC01 against the CIS Windows Server 2025 Benchmark. The individual checks show what the score is made of: password history, lockout thresholds, minimum password length, all evaluated by a specific command or registry key rather than asserted.




Figure 16.22: ADM01 against CIS Windows 11 Enterprise v3.0.0, the largest benchmark in the estate at 482 checks. This is the machine domain credentials are entered on, which is why its result matters as much as the domain controller's.




Figure 16.23: PVE01 against CIS Debian Linux 13. The failing checks here are recognisably different in character: filesystem partitioning and mount options rather than Group Policy, which is why the two families are not comparable by score.

Planned

Harden DC01 against a chosen set of failing checks, re-run the assessment, and record the before and after. That is the same shape as the management-plane isolation test in [docs/13](#) section 3.1, where a control was measured, changed, and measured again.

A score on its own proves nothing. **A score that moves proves the work happened**, and that is what an auditor is actually asking for.

A note on what Wazuh's vulnerability detection is, and is not

The 22 critical and 120 high findings on PVE01 are produced by matching the host's installed package inventory against CVE feeds. This is **credentialed, agent-based** detection: it knows precisely which versions are installed rather than inferring from service banners, so it is accurate and produces very few false positives.

It does not perform unauthenticated network scanning, so it will not tell you which services are exposed, which ports answer from another VLAN, whether TLS is weak, or whether default credentials are in use. Nor does it manage the remediation lifecycle: triage, ownership, deadlines, verification and reporting.

That distinction matters, because compliance regimes assess the process rather than the tool. **PCI-DSS 11.3** and **NIST SP 800-53 RA-5** both describe a cycle: identify, rank, remediate, verify, repeat. A scanner produces findings; vulnerability management is what is done with them. Extending this lab with a network scanner and a findings-aggregation layer is tracked in [docs/12](#).

11. What comes next, and where the rules go

Deploying the first agent is where the firewall model becomes concrete, and it is the point most people get backwards.

A pfSense rule governs who may start a conversation. Replies return through the state table and never need a rule. So the question for any service is not "who should reach this server?" but "who initiates?"

Wazuh agents connect **outbound** to the manager. The manager listens. Therefore:

To achieve this	The rule belongs on	Because
DC01 reports to Wazuh	ENTERPRISELAN (VLAN 50)	DC01 initiates, and it lives on VLAN 50
KALI01 reports to Wazuh	REDTEAM (VLAN 30)	KALI01 initiates
Administrator opens the dashboard	MANAGEMENT (VLAN 10)	the workstation initiates (MGMT-11 , built)

Almost nothing is added to the BLUETEAM tab itself. The instinct is to grant the SIEM broad reach on the reasoning that it must "see everything". That is backwards twice over: it does not grant what agents actually need, and it would give the machine holding all the security evidence a standing path across the entire estate.

The agent rules are `ENT-05` and `RED-04` , permitting TCP 1514 (events) and 1515 (enrollment) to `SIEM01_HOST` . They are registered in `docs/13` as they are built.

Also outstanding: Grafana and Prometheus are still installed on this host and should be removed once `MON01` exists, and Docker was removed as unused attack surface on a machine that has no need of it.

12. Standards alignment

Practice in this document	Standard
Centralised collection and retention of security events	NIST SP 800-53 AU-6 (Audit Record Review), AU-9 (Protection of Audit Information)
Continuous configuration assessment against a benchmark	NIST SP 800-53 CM-2, CM-6; CIS Controls v8, Control 4
Administrative access to the SIEM restricted and logged	NIST SP 800-53 AC-17, AU-2
Recovery of a failed service onto sound hardware, with placement justified	NIST SP 800-53 CP-10 (System Recovery)
Unnecessary software removed from a security-critical host	NIST SP 800-53 CM-7 (Least Functionality)

Practice in this document	Standard
Dashboard session bounded by an idle timeout, with the value justified	NIST SP 800-53 AC-11 (Session Lock), AC-12 (Session Termination)
Host identity verified out of band before trust was re-established	NIST SP 800-53 IA-3 (Device Identification and Authentication)
Role-based naming supporting asset inventory	NIST SP 800-53 CM-8
Segmentation between the monitored estate and the monitoring segment	NIST SP 800-53 SC-7, CIS Controls v8, Control 12

Related documents

- [docs/13-firewall-rulebase-governance.md](#) , the rule register including [MGMT-11](#) and the pending agent rules
- [docs/14-proxmox-storage-backup-capacity.md](#) , the naming convention and the role swap rationale
- [docs/15-kali-attack-host-build.md](#) , the same interface-direction reasoning applied to the attack segment
- [docs/02-security-monitoring.md](#) , the original Wazuh deployment on the failed Rocky Linux host